

# Hydra Infrastructure Management Guide

Student Container Platform Administration

Computer Science Department  
SUNY New Paltz

Last Updated: January 2025

## Contents

<b>1 System Overview</b>	<b>3</b>
1.1 Key Features . . . . .	3
1.2 Access URLs . . . . .	3
<b>2 System Architecture</b>	<b>3</b>
2.1 Architecture Diagram . . . . .	3
2.2 Component Overview . . . . .	4
2.3 Network Architecture . . . . .	4
<b>3 Authentication System</b>	<b>4</b>
3.1 SAML 2.0 SSO Flow . . . . .	4
3.2 Session Management . . . . .	4
<b>4 Container System</b>	<b>5</b>
4.1 Student Container Features . . . . .	5
4.2 Resource Limits . . . . .	5
4.3 Port Routing . . . . .	5
<b>5 File Structure</b>	<b>6</b>
<b>6 Common Operations</b>	<b>6</b>
6.1 View Running Containers . . . . .	6
6.2 Access Container Shell . . . . .	6
6.3 View Container Logs . . . . .	6
6.4 Restart a Container . . . . .	6
6.5 Remove a Stuck Container . . . . .	6
6.6 Build Student Container Images . . . . .	7
<b>7 Service Management</b>	<b>7</b>
7.1 Restart Main Services . . . . .	7
7.2 Build and Deploy . . . . .	7
7.3 View Service Logs . . . . .	7
7.4 Check Traffic Routing . . . . .	7
<b>8 Troubleshooting</b>	<b>7</b>
8.1 Authentication Issues . . . . .	7
8.2 Container Issues . . . . .	8
8.3 Service-Specific Issues . . . . .	8

<b>9 Backup and Recovery</b>	<b>8</b>
9.1 Databas Backup . . . . .	8
9.2 Volum Backup . . . . .	8
<b>10 Environment Configuration</b>	<b>9</b>
10.1 R quir d Variabl s . . . . .	9
10.2 Optional Variabl s . . . . .	9
<b>11 References</b>	<b>9</b>

## 1 System Overview

Hydra is a containerized development platform providing persistent development environments for Computer Science students and faculty at SUNY New Paltz. The system uses SAML 2.0 Single Sign-On via Azure AD and Docker for container orchestration.

### 1.1 Key Features

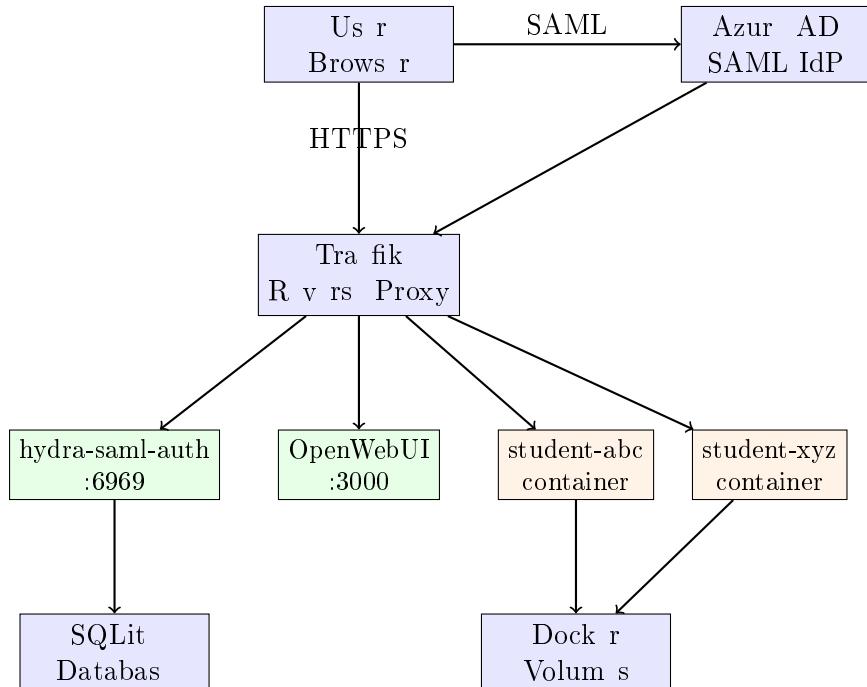
- SSO Authentication:** Azure AD SAML 2.0 with automatic user provisioning
- Persistent Containers:** One development environment per student with data persistence
- Built-in Services:** VS Code (code-server), Jupyter Notebook, Docker-in-Docker
- Dynamic Routing:** Traffic-based routing for custom web applications
- Resource Management:** Per-container CPU and memory limits
- Integration:** OpenWebUI (GPT) and n8n account management

### 1.2 Access URLs

Service	URL	Description
Dashboard	<a href="https://hydr.newpaltz.edu/dashboard">https://hydr.newpaltz.edu/dashboard</a>	Main user interface
OpenWebUI	<a href="https://gpt.hydr.newpaltz.edu/">https://gpt.hydr.newpaltz.edu/</a>	AI chat interface
VS Code	<a href="https://hydr.newpaltz.edu/students/{user}/vscode">https://hydr.newpaltz.edu/students/{user}/vscode</a>	Browser IDE
Jupyter	<a href="https://hydr.newpaltz.edu/students/{user}/jupyter">https://hydr.newpaltz.edu/students/{user}/jupyter</a>	Notebooks

## 2 System Architecture

### 2.1 Architecture Diagram



## 2.2 Component Overview

Component	Port	Description
Tra fik	80, 443	Revers proxy, TLS termination, routing
hydra-saml-auth	6969	SAML auth, dashboard, container management
OpenW bUI	3000	AI chat interface (Ollama frontend)
Student Containers	Dynamic	Prisoner database environments

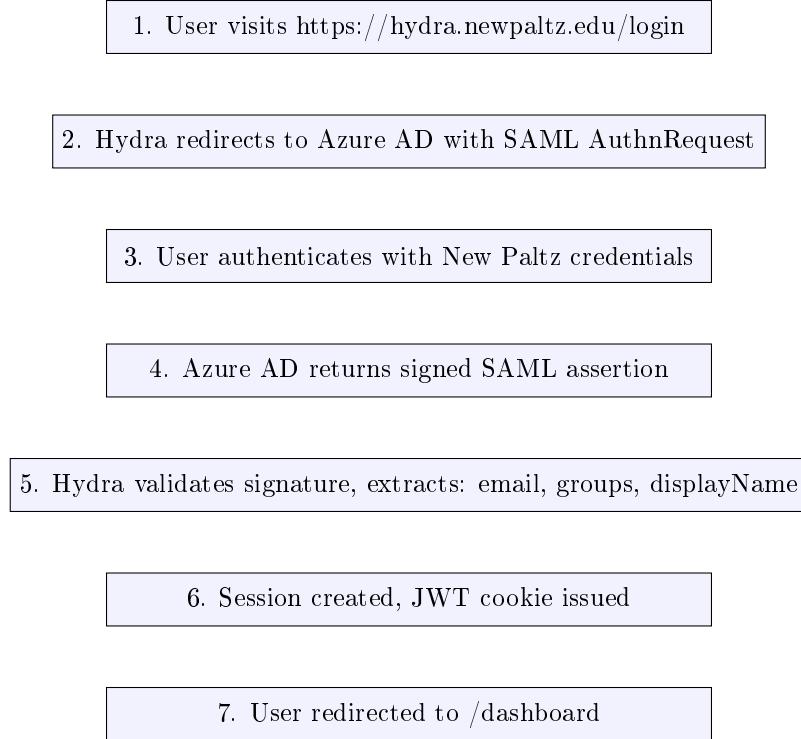
## 2.3 Network Architecture

Students containers operate on an isolated Docker network (`hydr_students_net`) with:

- No direct internal access (configurable)
- Internal DNS resolution
- Traffic-matching external access via ForwardAuth

## 3 Authentication System

### 3.1 SAML 2.0 SSO Flow



### 3.2 Session Management

Sessions are managed via:

- Express Session:** Server-side session storage in SQLit
- JWT Cookie:** Site-wide authentication cookie for cross-service SSO

- **JWKS Endpoint:** Public key endpoint for JWT verification by other services

#### JWT Configuration:

- TTL: Configurable via `JWT_TTL_SECONDS` (default: 86400)
- Algorithm: RS256
- Cookie Domain: `.newplatformz.edu`

## 4 Container System

### 4.1 Student Container Features

Each student container has a single persistent container with:

Feature	Details
Node.js	Latest LTS via nvm
Python	3.11+ with pip, venv, Jupyter
Java	OpenJDK 21
Docker	Full Docker-in-Docker support (privileged mode)
VS Code	Code-server browser IDE
Jupyter	Notebook and JupyterLab
Tools	Git, curl, wget, build-utils, etc.

### 4.2 Resource Limits

Resource	Limit
RAM	4GB per container
CPU	2 cores per container
Storage	Unlimited (host disk limit)

**Security Note:** Students run in privileged mode to support Docker-in-Docker. This grants elevated access. Monitor for abuse and consider disabling for untrusted users.

### 4.3 Port Routing

Students can expose web applications through custom routes:

- Default routes: `/students/{username}/vscode`, `/students/{username}/jupyter`
- Custom routes added via dashboard UI
- Reserved ports: 8443 (code-server), 8888 (Jupyter)
- All routes protected by ForwardAuth

## 5 File Structure

```
hydra-saml-auth/
|-- index.js                      # Main entry: SAML, JWT/JWKS, routes,
  WebSocket
|-- db.js                          # SQLite database initialization
|-- routes/
|   |-- containers.js             # Container lifecycle, services, ports, logs
|   |-- webui-api.js              # OpenWebUI account proxy
|   |-- n8n-api.js                # n8n account management
|   |-- servers-api.js            # Cluster status endpoints
|   |-- admin.js                  # Admin panel routes
|-- services/
|   |-- activity-logger.js        # Activity tracking
|   |-- email-notifications.js   # Email alerts
|-- views/                         # EJS templates
|-- student-container/
|   |-- Dockerfile               # Ubuntu 22.04 + dev tools
|   |-- supervisord.conf          # Process manager config
|   |-- entrypoint.sh             # Container startup
|-- docker-compose.yaml            # Production stack
|-- docs/                          # Documentation
```

## 6 Common Operations

### 6.1 View Running Containers

```
docker ps --filter "name=student -"
```

### 6.2 Access Container Shell

```
docker exec -it student-<username> /bin/bash
```

### 6.3 View Container Logs

```
docker logs -f student-<username> --tail=100
```

### 6.4 Restart a Container

```
docker restart student-<username>
```

### 6.5 Remove a Stuck Container

```
docker rm -f student-<username>
```

## 6.6 Rebuild Student Container Image

```
cd student-container  
docker build -t hydra-student-container:latest .
```

**Note:** Students with existing containers must recreate them to use updated images.

# 7 Service Management

## 7.1 Restart Main Service

```
docker compose restart hydra-saml-auth
```

## 7.2 Rebuild and Redeploy

```
docker compose build hydra-saml-auth  
docker compose up -d hydra-saml-auth
```

## 7.3 View Service Logs

```
docker compose logs -f hydra-saml-auth
```

## 7.4 Check Traefik Routing

```
docker compose logs traefik | grep -i error  
curl -I https://hydra.newpaltz.edu/
```

# 8 Troubleshooting

## 8.1 Authentication Issues

Symptom	Solution
SAML assertion invalid	Verify METADATA_URL and SAML_SP_ENTITY_ID match Azure config exactly
Cookies not set	Check COOKIE_DOMAIN, ensure HTTPS, check browser settings
JWT verification fails	Verify JWKS endpoint accessible, check key rotation

## 8.2 Container Issues

Symptom	Solution
Container won't initialize	Verify <code>hydr -student-container:1</code> test image exists
Container 404	Check container is on <code>hydr_students_net</code> , traffic running
Service won't start	Check supervisor logs inside container
Port routing fails	Verify port not forwarded (8443, 8888) and not in use

## 8.3 Service-Specific Issues

- VS Code not loading:** Check code-server process, ForwardAuth working
- Jupyter issues:** Verify `NotebookApp.base_url` setting
- Docker-in-Docker fails:** Container must have privileged mode
- Files not persisting:** Only `/home/student/` is persisted

# 9 Backup and Recovery

## 9.1 Database Backup

```
# Backup SQLite database
sqlite3 /app/data/webui.db ".backup '/backups/hydra-$(date +\%Y\%m\%d).db"
,"

# Automated daily backup (add to crontab)
0 2 * * * sqlite3 /app/data/webui.db ".backup '/backups/hydra-$(date
+\%Y\%m\%d).db,"
```

## 9.2 Volume Backup

```
# List student volumes
docker volume ls | grep student-

# Backup a volume
docker run --rm -v student-<user>-data:/data -v $(pwd):/backup \
alpine tar cvf /backup/student-<user>-backup.tar /data
```

## 10 Environment Configuration

### 10.1 Required Variables

Variable	Description
BASE_URL	External URL (https://hydra.nwpaltz.du)
METADATA_URL	Azure AD federation metadata URL
SAML_SP_ENTITY_ID	SP Entity ID (must match Azure exactly)
COOKIE_DOMAIN	Cookie scope (.nwpaltz.du)
PORT	Service port (default: 6969)
DB_PATH	Database path (/app/data/wbui.db)

### 10.2 Optional Variables

Variable	Description
PUBLIC_STUDENTS_BASE	Student URL base (https://hydra.nwpaltz.du/students)
JWT_TTL_SECONDS	JWT token lifetime
JWT_PRIVATE_KEY_FILE	Path to JWT signing key
JWT_PUBLIC_KEY_FILE	Path to JWT verification key

## 11 References

- Docker Documentation: <https://docs.docker.com/>
- Traefik Documentation: <https://doc.traefik.io/traefik/>
- SAML 2.0 Specification: <https://docs.oasis-open.org/security/saml/v2.0/>
- Azure AD SAML: <https://docs.microsoft.com/en-us/azure/active-directory/develop/single-sign-on-single-protocol>
- Code-server: <https://coder.com/docs/code-server/1-test>
- Jupyter: <https://jupyter.org/documentation>