

Hydra Installation & Setup Guide

For System Administrators

Computer Science Department
SUNY New Paltz

Last Updated: January 2025

Contents

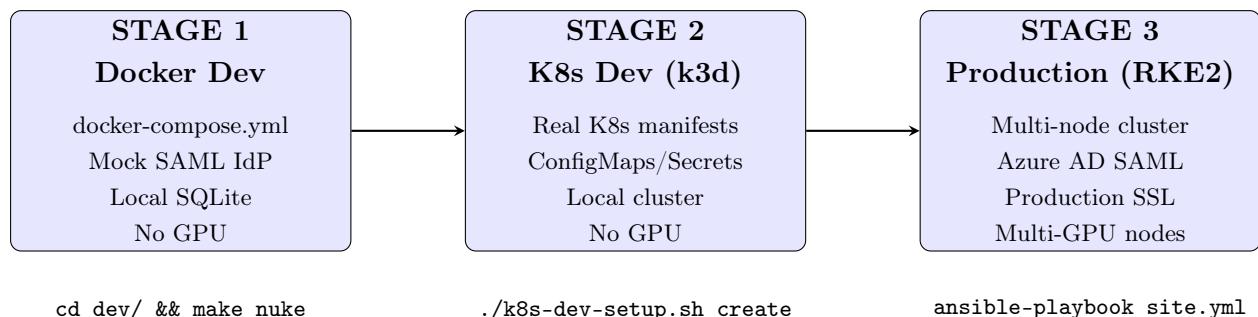
1 Deployment Overview	2
1.1 Deployment Pipeline	2
1.2 Quick Start Commands	2
1.3 Multi-Node Production Architecture	2
2 Prerequisites	3
2.1 System Requirements	3
2.2 Required Access	3
2.3 Software Dependencies	3
3 Installation Steps	3
3.1 Step 1: Clone Repository	3
3.2 Step 2: Build Student Container Image	3
3.3 Step 3: Configure Environment	4
3.4 Step 4: Generate JWT Keys	4
3.5 Step 5: Configure Azure AD	5
3.6 Step 6: Configure Traefik	5
3.7 Step 7: Start Services	5
3.8 Step 8: Verify Installation	6
4 Post-Installation Setup	6
4.1 Create Admin User	6
4.2 Test Container Creation	6
4.3 Configure Backup Cron	6
5 Network Architecture	7
5.1 Deployment Diagram	7
5.2 Port Mapping	7
6 Security Configuration	7
6.1 TLS Certificates	7
6.2 Container Security	8
7 Useful Aliases	8
8 Monitoring	9
8.1 Log Locations	9
8.2 Health Endpoints	9

9 Troubleshooting	10
9.1 Common Issues	10
9.2 Debug Commands	10
10 Upgrade Procedures	10
10.1 Updating the Application	10
10.2 Updating Student Container Image	10
11 References	11
11.1 Core Infrastructure	11
11.2 Authentication & Security	11
11.3 Development Tools	11
11.4 AI & Automation Services	11
11.5 Kubernetes (Future Migration)	12
11.6 Storage & Backup	12
11.7 Game Servers	12
11.8 Monitoring & Logging	12
11.9 SUNY New Paltz Resources	12

1 Deployment Overview

This guide covers three deployment stages, from local development to production:

1.1 Deployment Pipeline



1.2 Quick Start Commands

```

# STAGE 1: Docker Development (start here)
cd dev/
make nuke          # Complete reset + rebuild

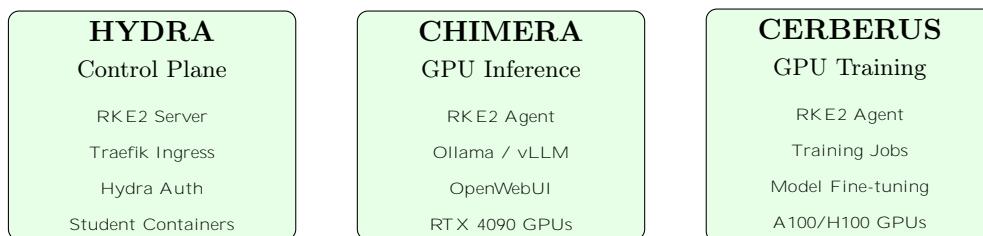
# STAGE 2: Kubernetes Development (after Docker works)
./k8s-dev-setup.sh create

# STAGE 3: Production (after K8s validation)
cd ansible/
ansible-playbook -i inventory.yml playbooks/site.yml

```

Recommended Workflow: Always validate changes in Stage 1 (Docker) before promoting to Stage 2 (k3d), and validate in Stage 2 before deploying to Stage 3 (Production).

1.3 Multi-Node Production Architecture



2 Prerequisites

2.1 System Requirements

Component	Requirement
Operating System	Ubuntu 22.04 LTS (recommended)
Docker	24.0+ with Compose V2
RAM	Minimum 8GB (16GB+ recommended)
Storage	100GB+ SSD
Network	Static IP, ports 80/443 open

2.2 Required Access

- Azure AD admin access for SAML Enterprise Application setup
- DNS control for `hydra.newpaltz.edu` and subdomains
- TLS certificate (Let's Encrypt or institutional cert)
- SSH access to the host server

2.3 Software Dependencies

```
# Install Docker
curl -fsSL https://get.docker.com | sh
sudo usermod -aG docker $USER

# Verify Docker Compose V2
docker compose version
# Should show: Docker Compose version v2.x.x

# Install useful tools
sudo apt install -y git curl jq sqlite3
```

3 Installation Steps

3.1 Step 1: Clone Repository

```
git clone https://github.com/your-org/hydra-saml-auth.git
cd hydra-saml-auth
```

3.2 Step 2: Build Student Container Image

```
cd student-container
docker build -t hydra-student-container:latest .
cd ..
```

This image includes: Ubuntu 22.04, Node.js, Python 3.11+, Java 21, Docker-in-Docker, code-server, and Jupyter.

3.3 Step 3: Configure Environment

Create .env file in the project root:

```
# === Core Settings ===
PORT=6969
BASE_URL=https://hydra.newpaltz.edu
COOKIE_DOMAIN=.newpaltz.edu

# === SAML Configuration ===
METADATA_URL=https://login.microsoftonline.com/YOUR_TENANT/
    federationmetadata/2007-06/federationmetadata.xml
SAML_SP_ENTITY_ID=hydra-auth
SAML_CALLBACK_URL=https://hydra.newpaltz.edu/auth/callback

# === Database ===
DB_PATH=/app/data/webui.db

# === JWT Settings ===
JWT_TTL_SECONDS=86400
JWT_KEY_ID=hydra-key-1
JWT_PRIVATE_KEY_FILE=/app/certs/jwt-private.pem
JWT_PUBLIC_KEY_FILE=/app/certs/jwt-public.pem

# === Student Containers ===
PUBLIC_STUDENTS_BASE=https://hydra.newpaltz.edu/students

# === OpenWebUI Integration ===
OPENWEBUI_URL=https://gpt.hydra.newpaltz.edu
OPENWEBUI_API_KEY=your-openwebui-api-key

# === n8n Workflow Automation ===
N8N_URL=https://n8n.hydra.newpaltz.edu
N8N_WEBHOOK_URL=https://n8n.hydra.newpaltz.edu/webhook
```

3.4 Step 4: Generate JWT Keys

```
mkdir -p certs
openssl genrsa -out certs/jwt-private.pem 2048
openssl rsa -in certs/jwt-private.pem -pubout -out certs/jwt-public.pem
```

3.5 Step 5: Configure Azure AD

1. Go to Azure Portal > Azure Active Directory > Enterprise Applications

2. Click "New application" > "Create your own application"

3. Name: "Hydra Auth", select "Non-gallery application"

4. Go to Single sign-on > SAML

5. Set Identifier (Entity ID): hydra-auth

6. Set Reply URL: https://hydra.newpaltz.edu/auth/callback

7. Download Federation Metadata XML, note the URL

8. Assign users/groups who should have access

Critical: The Entity ID in Azure must **exactly match** SAML_SP_ENTITY_ID in your .env file.

3.6 Step 6: Configure Traefik

Ensure docker-compose.yaml has proper Traefik configuration:

```
# Key Traefik labels for hydra-saml-auth service:  
labels:  
  - "traefik.enable=true"  
  - "traefik.http.routers.hydra.rule=Host('hydra.newpaltz.edu')"  
  - "traefik.http.routers.hydra.entrypoints=websecure"  
  - "traefik.http.routers.hydra.tls=true"  
  - "traefik.http.services.hydra.loadbalancer.server.port=6969"
```

3.7 Step 7: Start Services

```
# Build and start all services  
docker compose build  
docker compose up -d  
  
# Verify services are running  
docker compose ps  
  
# Check logs  
docker compose logs -f hydra-saml-auth
```

3.8 Step 8: Verify Installation

```
# Test HTTPS access
curl -I https://hydra.newpaltz.edu/

# Should redirect to Azure AD login (302 to login.microsoftonline.com)

# Check JWKS endpoint
curl https://hydra.newpaltz.edu/.well-known/jwks.json

# Test OpenWebUI
curl -I https://gpt.hydra.newpaltz.edu/

# Test n8n (if configured)
curl -I https://n8n.hydra.newpaltz.edu/
```

4 Post-Installation Setup

4.1 Create Admin User

After first login via SAML, promote a user to admin:

```
# Access database
sqlite3 /app/data/webui.db

# Find user ID
SELECT id, email, role FROM users WHERE email LIKE '%admin%';

# Set as admin
UPDATE users SET role = 'admin' WHERE email = 'admin@newpaltz.edu';
```

4.2 Test Container Creation

1. Log in to <https://hydra.newpaltz.edu/dashboard>
2. Navigate to "Containers" tab
3. Click "Initialize Container"
4. Verify VS Code and Jupyter are accessible at:
 - <https://hydra.newpaltz.edu/students/{username}/vscode>
 - <https://hydra.newpaltz.edu/students/{username}/jupyter>

4.3 Configure Backup Cron

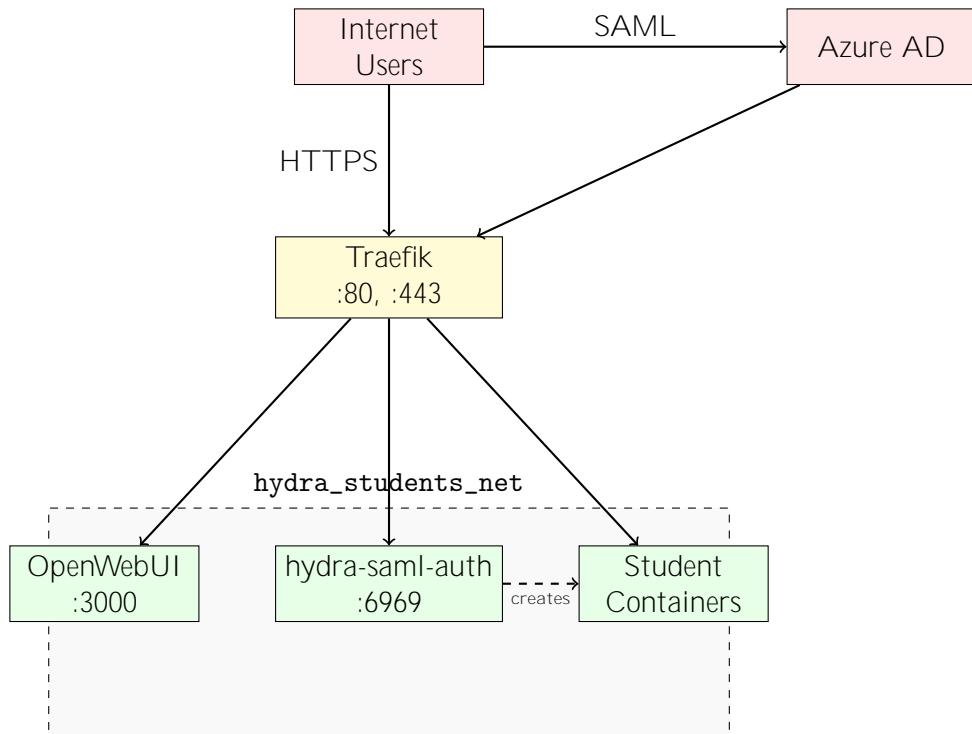
```
# Add to root crontab
crontab -e

# Daily database backup at 2 AM
0 2 * * * docker exec hydra-saml-auth sqlite3 /app/data/webui.db ". 
    backup '/backups/hydra-$(date +\%Y\%m\%d).db'"

# Weekly cleanup of old backups
0 3 * * 0 find /backups -name "hydra-*.\.db" -mtime +30 -delete
```

5 Network Architecture

5.1 Deployment Diagram



5.2 Port Mapping

Service	Internal	External	Protocol
Traefik	-	80, 443	HTTP/HTTPS
hydra-saml-auth	6969	via Traefik	HTTP
OpenWebUI	3000	via Traefik	HTTP
Student VS Code	8443	/students/{user}/vscode	HTTP
Student Jupyter	8888	/students/{user}/jupyter	HTTP

6 Security Configuration

6.1 TLS Certificates

For Let's Encrypt (recommended):

```

# Traefik handles automatic certificate management
# Add to docker-compose.yaml traefik command:
command:
  - "--certificatesresolvers.letsencrypt.acme.email=admin@newpaltz.edu"
  - "--certificatesresolvers.letsencrypt.acme.storage=/letsencrypt/acme
    .json"
  - "--certificatesresolvers.letsencrypt.acme.httpchallenge.entrypoint=
    web"

```

For institutional certificates:

```

# Mount certificates in Traefik
volumes:

```

```
- ./certs/cert.pem:/certs/cert.pem:ro  
- ./certs/key.pem:/certs/key.pem:ro  
  
# Configure in dynamic config  
tls:  
    certificates:  
        - certFile: /certs/cert.pem  
          keyFile: /certs/key.pem
```

6.2 Container Security

Student containers run in **privileged mode** by default to enable Docker-in-Docker. This grants significant access. Consider:

- Restricting to trusted users only
- Monitoring container activity
- Implementing resource quotas

7 Useful Aliases

Add to `~/.bashrc`:

```
# ===== HYDRA MANAGEMENT ALIASES =====  
  
# Docker shortcuts  
alias dps='docker ps --format "table {{.Names}}\t{{.Status}}\t{{.Ports}}"'  
alias dlogs='docker compose logs -f'  
alias dexec='docker exec -it'  
  
# Hydra-specific  
alias hydra-logs='docker compose logs -f hydra-saml-auth'  
alias hydra-restart='docker compose restart hydra-saml-auth'  
alias hydra-rebuild='docker compose build hydra-saml-auth && docker compose up -d hydra-saml-auth'  
  
# Student container management  
alias students='docker ps --filter "name=student-" --format "table {{.Names}}\t{{.Status}}"'  
alias student-logs='docker logs -f'  
alias student-shell='docker exec -it'  
alias student-stop='docker stop'  
alias student-rm='docker rm -f'  
  
# Find student by partial name  
findstudent() {  
    docker ps --filter "name=student-" --format "{{.Names}}" | grep -i "$1"  
}  
  
# Get student container stats  
student-stats() {  
    docker stats --no-stream --filter "name=student-"  
}
```

```
# Backup database
backup-db() {
    docker exec hydra-saml-auth sqlite3 /app/data/webui.db ".backup '/
        backups/hydra-$(date +%Y%m%d-%H%M%S).db'"
    echo "Backup created: hydra-$(date +%Y%m%d-%H%M%S).db"
}

# Quick health check
hydra-health() {
    echo "==== Services ==="
    docker compose ps
    echo ""
    echo "==== Student Containers ==="
    docker ps --filter "name=student-" --format "table {{.Names}}\t{{.
        Status}}\t{{.RunningFor}}"
    echo ""
    echo "==== Resource Usage ==="
    docker stats --no-stream --format "table {{.Name}}\t{{.CPUPerc}}\t
        {{.MemUsage}}" | head -10
}
```

Apply changes:

```
source ~/.bashrc
```

8 Monitoring

8.1 Log Locations

Component	Command
Main service	docker compose logs hydra-saml-auth
Traefik	docker compose logs traefik
Student container	docker logs student-<username>
All services	docker compose logs -f

8.2 Health Endpoints

```
# Check main service
curl https://hydra.newpaltz.edu/health

# Check JWKS (JWT verification)
curl https://hydra.newpaltz.edu/.well-known/jwks.json

# Check Traefik dashboard (if enabled)
curl http://localhost:8080/api/overview

# Check OpenWebUI health
curl https://gpt.hydra.newpaltz.edu/health
```

9 Troubleshooting

9.1 Common Issues

Issue	Solution
SAML login fails	Verify <code>METADATA_URL</code> accessible, Entity ID matches exactly
Student container 404	Check container on <code>hydra_students_net</code> , Traefik running
Container won't start	Verify <code>hydra-student-container:latest</code> image built
Permission denied	Check Docker socket permissions, user in docker group
Database locked	Restart service, check for multiple writers

9.2 Debug Commands

```
# Check Docker networks
docker network ls
docker network inspect hydra_students_net

# Verify image exists
docker images | grep hydra-student-container

# Check container networking
docker exec student-<user> curl -I http://hydra-saml-auth:6969/

# View Traefik routes
docker exec traefik cat /etc/traefik/traefik.yml
```

10 Upgrade Procedures

10.1 Updating the Application

```
# Pull latest code
git pull origin main

# Rebuild and restart
docker compose build hydra-saml-auth
docker compose up -d hydra-saml-auth

# Verify
docker compose logs -f hydra-saml-auth
```

10.2 Updating Student Container Image

```
cd student-container
docker build -t hydra-student-container:latest .
```

Existing student containers continue using the old image. Students must delete and recreate their container to get the updated image.

11 References

11.1 Core Infrastructure

- Docker Documentation: <https://docs.docker.com/>
- Docker Compose: <https://docs.docker.com/compose/>
- Docker-in-Docker (DinD): <https://docs.docker.com/engine/reference/run/#runtime-privilege-and-sandbox>
- Traefik Reverse Proxy: <https://doc.traefik.io/traefik/>
- Traefik ForwardAuth: <https://doc.traefik.io/traefik/middlewares/http/forwardauth/>
- Let's Encrypt (TLS Certificates): <https://letsencrypt.org/docs/>

11.2 Authentication & Security

- SAML 2.0 Specification: <https://docs.oasis-open.org/security/saml/v2.0/>
- Azure AD SAML Setup: <https://learn.microsoft.com/en-us/azure/active-directory/develop/single-sign-on-saml-protocol>
- Azure AD Enterprise Applications: [https://learn.microsoft.com/en-us/azure/active-directory/manage-apps/](https://learn.microsoft.com/en-us/azure/active-directory/manage-apps)
- passport-saml (Node.js): <https://github.com/node-saml/passport-saml>
- JSON Web Tokens (JWT): <https://jwt.io/introduction>
- JWKS (JSON Web Key Sets): <https://auth0.com/docs/secure/tokens/json-web-tokens/json-web-key-sets>

11.3 Development Tools

- code-server (VS Code in Browser): <https://coder.com/docs/code-server/latest>
- Jupyter Notebook: <https://jupyter.org/documentation>
- JupyterLab: <https://jupyterlab.readthedocs.io/en/stable/>
- Node.js: <https://nodejs.org/en/docs/>
- Python: <https://docs.python.org/3/>
- OpenJDK: <https://openjdk.org/>

11.4 AI & Automation Services

- OpenWebUI (GPT Interface): <https://docs.openwebui.com/>
- Ollama (Local LLM Runtime): <https://ollama.ai/>
- n8n Workflow Automation: <https://docs.n8n.io/>
- n8n Self-Hosting: <https://docs.n8n.io/hosting/>

11.5 Kubernetes (Future Migration)

- RKE2 (Rancher Kubernetes Engine): <https://docs.rke2.io/>
- Kubernetes Documentation: <https://kubernetes.io/docs/>
- NVIDIA GPU Operator: <https://docs.nvidia.com/datacenter/cloud-native/gpu-operator/>
- Longhorn (Distributed Storage): <https://longhorn.io/docs/>

11.6 Storage & Backup

- ZFS on Linux: <https://openzfs.github.io/openzfs-docs/>
- RAID Levels Explained: https://en.wikipedia.org/wiki/Standard_RAID_levels
- RAID 10 Configuration: <https://www.prepressure.com/library/technology/raid>
- SQLite Documentation: <https://www.sqlite.org/docs.html>
- Docker Volumes: <https://docs.docker.com/storage/volumes/>

11.7 Game Servers

- Minecraft Server (Java Edition): https://minecraft.wiki/w/Tutorials/Setting_up_a_server
- Minecraft EULA: <https://www.minecraft.net/en-us/eula>
- itzg/minecraft-server (Docker Image): <https://docker-minecraft-server.readthedocs.io/>

11.8 Monitoring & Logging

- Docker Logs: <https://docs.docker.com/config/containers/logging/>
- Supervisord (Process Manager): <http://supervisord.org/>
- htop (Process Viewer): <https://htop.dev/>

11.9 SUNY New Paltz Resources

- Hydra Dashboard: <https://hydra.newpaltz.edu/dashboard>
- OpenWebUI (GPT): <https://gpt.hydra.newpaltz.edu/>
- SUNY New Paltz ITS: <https://www.newpaltz.edu/its/>

Installation Complete!

After completing these steps:

1. Visit <https://hydra.newpaltz.edu/login>
2. Authenticate via Azure AD with your SUNY New Paltz credentials
3. Navigate to the dashboard
4. Create a test container to verify the setup
5. Access VS Code at /students/{username}/vscode

6. Access Jupyter at `/students/{username}/jupyter`

For ongoing management, refer to the **Hydra Infrastructure Management Guide**.