

# Hydra RKE2 Cluster Guide

Storage Architecture & GPU Resource Management

Computer Science Department  
SUNY New Paltz

February 2026

## Contents

<b>1</b>	<b>Overview</b>	<b>2</b>
<b>2</b>	<b>Cluster Topology</b>	<b>2</b>
2.1	GPU Resources . . . . .	2
<b>3</b>	<b>Kubernetes Namespaces</b>	<b>2</b>
<b>4</b>	<b>Storage Architecture</b>	<b>2</b>
4.1	Ti rd Storage Design . . . . .	2
4.2	ZFS Configuration (Hydra) . . . . .	3
4.3	Storage Classes . . . . .	3
<b>5</b>	<b>Node Labeling &amp; Scheduling</b>	<b>3</b>
5.1	Node Labels . . . . .	3
5.2	GPU Tolerations . . . . .	3
<b>6</b>	<b>Resource Presets</b>	<b>4</b>
<b>7</b>	<b>Manifest Structure</b>	<b>4</b>
<b>8</b>	<b>Deployment Procedures</b>	<b>4</b>
8.1	Initial Setup . . . . .	4
8.2	Health Check . . . . .	5
<b>9</b>	<b>Backup Strategy</b>	<b>5</b>
<b>10</b>	<b>Migration Status</b>	<b>5</b>

## 1 Overview

The Hydra cluster is a 3-node RKE2 Kubernetes deployment serving the SUNY New Paltz Computer Science department. It provides a containerized environment for students and GPU resources for machine learning workloads.

**Repository:** <https://github.com/compsci-suny-newp-ltz/hydra-k3s>

**Orchestration:** Rancher RKE2 v1.28

**Container Runtime:** containerd 1.7.7

## 2 Cluster Topology

Node	IP	Role	RAM	Storage
Hydra	192.168.1.160	Control plane, student containers	256 GB	21 TB ZFS RAID-10
Chimera	192.168.1.150	GPU infrastructure (OpenWarpUI + Ollama)	251 GB	3.5 TB NVM
Crabfish	192.168.1.233	GPU training	64 GB	3.6 TB NVM

Table 1: Node hardware specifications

### 2.1 GPU Resources

Node	GPU Model	Count	VRAM	Use Case
Chimera	RTX 3090	3	72 GB	Ollama inference, concurrent users
Crabfish	RTX 5090	2	64 GB	Student training, research

Table 2: GPU allocation across the cluster (136 GB total VRAM)

**GPU Scheduling:** Chimera GPUs are shared with OpenWarpUI/Ollama. For dedicated GPU access, student workloads should target Crabfish via the `gpu_training` resource provider.

## 3 Kubernetes Namespaces

Namespace	Purpose
hydra-system	Core services: Traefik, hydra-auth, cs-lab-backend, cs-lab-db
hydra-students	Student workload pods (on primary student)
gpu-operator	NVIDIA GPU operator and driver plugin
kube-system	RKE2 system components

Table 3: Namespace organization

## 4 Storage Architecture

### 4.1 Tiered Storage Design

The cluster uses a 3-tier storage model:

1. **Tier 1 — Fast NVMe Cache (8.8 TB):** Chimera and Cerberus NVM drives for hot data, active models, and training workspaces.
2. **Tier 2 — Primary Storage (21 TB):** Hydra ZFS RAID-10 array for student content, model repositories, and system backups.
3. **Tier 3 — Archive:** Long-term storage for graduated student data and completed projects.

## 4.2 ZFS Configuration (Hydra)

```
# RAID-10 array: 6 drives, 21TB usable
/dev/md0: active raid10 sdc sdd sdf sde sdb sda
      22500996096 blocks [6/6] [UUUUUU]

# Mount point
/dev/md0 21T 54G 20T 1% /data
```

## 4.3 Storage Classes

StorageClass	Backend	Use Case
hydr -loc	ZFS on Hydra	Student PVCs, system data
hydr -nfs	NFS export from Hydra	Cross-node access (GPU nodes)
hydr -hot	NVM on GPU nodes	Active training data

Table 4: Kubernetes Storage Class definitions

## 5 Node Labeling & Scheduling

Pods are scheduled to appropriate nodes using labels and tolerations.

### 5.1 Node Labels

```
# Hydra (control plane)
hydra.node-role=control-plane

# Chimera (inference)
hydra.node-role=inference
hydra.gpu-enabled=true

# Cerberus (training)
hydra.node-role=training
hydra.gpu-enabled=true
```

### 5.2 GPU Tolerations

GPU workloads must include the NVIDIA tolerance:

```
1 tolerations:
2   - key: nvidia.com/gpu
3     operator: Exists
4     effect: NoSchedule
```

## 6 Resource Presets

Students select resource presets when requesting storage containers. These map to Kubernetes storage resources and limits.

Preset	Memory	CPUs	Storage	GPU	Auto-Approve
Minimal	1 GB	0.5	5 GB	—	Yes
Conservative	1.5 GB	1	10 GB	—	Yes
Standard	2 GB	1	20 GB	—	Yes
Enhanced	4 GB	2	40 GB	—	Yes
GPU Infrastructure	32 GB	8	100 GB	1	No
GPU Training	48 GB	16	200 GB	2	No

Table 5: Resource presets defined in `config/resources.js`

**GPU presets require admin approval.** GPU Infrastructure targets Chimera; GPU Training targets Cerberus. Capacity is limited — approved based on current utilization.

## 7 Manifest Structure

```
hydra-k3s/
  manifests/
    cache-config/          # NVMe cache tier configuration
    model-cache/           # Model caching policies
    monitoring/            # Prometheus, alerting
    storage-controller/   # Tiered storage controller
  scripts/
    setup-hydra-storage.sh
    setup-chimera-cache.sh
    setup-cerberus-workspace.sh
    deploy-controllers.sh
    migrate-data.sh
  docs/
    phase1-storage.md
    phase2-fast-tier.md
    phase3-controllers.md
    phase4-migration.md
```

## 8 Deployment Procedures

### 8.1 Initial Setup

```
# Phase 1: Setup storage (Day 1)
./scripts/setup-hydra-storage.sh

# Phase 2: Configure fast tier (Day 2)
kubectl apply -f manifests/cache-config/

# Phase 3: Deploy controllers (Day 3)
kubectl apply -f manifests/storage-controller/
kubectl apply -f manifests/monitoring/
```

```
# Phase 4: Migrate data (Days 4-5)
./scripts/migrate-data.sh
```

## 8.2 Health Check

```
export KUBECONFIG=/etc/rancher/rke2/rke2.yaml

# Node status
kubectl get nodes -o wide

# Problem pods
kubectl get pods -A | grep -v Running | grep -v Completed

# GPU availability
kubectl get nodes -o custom-columns=\
"NAME:.metadata.name,GPU:.status.capacity.nvidia\.com/gpu"

# Student pods
kubectl get pods -n hydra-students --sort-by=.metadata.name

# RAID health
cat /proc/mdstat | head -5
```

## 9 Backup Strategy

- **Weekly OS backups:** backup-cluster.sh via cron (rsync to /mnt/sdh4/backups)
- **ZFS snapshots:** Hourly for active datasets, daily for archives
- **etcd snapshots:** Automatic via RKE2 (stored on Hydra sdb)
- **Database dumps:** MariaDB and PostgreSQL included in weekly backup

## 10 Migration Status

Phase	Task	Status
Phase 1	Storage pool created	Completed
Phase 2	Fast timer configured	Pending
Phase 3	Controllers deployed	Pending
Phase 4	Data migration	Pending

Table 6: Docker-to-Kubernetes migration progress